

Phase 4 Formal Framework: Decoupled Planning and Learning Burdens in Symbolic Action Models

MSc Thesis Research Program
Target Venue: ICAPS 2027 / AAAI 2027

October 1, 2026

Abstract

This document establishes the formal theoretical framework, research objectives, research questions, hypotheses, algorithms, and statistical analysis protocol for Phase 4 of the research program: *Mechanic-Induced Planning-Model Burden Decomposition*. We formalize the distinction between Planning Burden (H_P) and Bounded Learning Burden (B_{bounded}), prove the Decoupling Theorem and Zone B divergence (the critical blindspot where models appear accurate under passive evaluation yet induce catastrophic planning failure), present complete algorithmic pseudocode, specify an audited 8-domain International Planning Competition (IPC) empirical matrix, and detail the statistical testing and compute-budgeting pipeline.

1 Context, Problem Statement, and Research Gaps

1.1 Discrete Symbolic Planning vs. Continuous Control World Models

In continuous deep world models (e.g., DreamerV3, PlaNet), state representations lie in metric spaces ($\mathbb{R}^d$), and dynamics prediction errors typically manifest as gradual drift measured by mean squared error (MSE) across rollout horizons. In contrast, **symbolic planning environments possess discrete, combinatorial topology governed by first-order logic constraints**.

A single omitted precondition in an action schema does not cause a continuous numerical perturbation; rather, it performs a discrete topological relaxation on the state-space reachability graph $\mathcal{R}(s_0)$. This creates *phantom transition shortcuts* that a forward search planner greedily exploits, causing catastrophic real-world execution failure.

1.2 The Three-Tier Problem Statement

1. **Tier 1 (Known Phenomenon):** High passive transition prediction accuracy ($A_{\text{pred}} \geq 98\%$) does not guarantee planning success (*Verified-vs-Correct Gap*, Aguilar Martín et al., 2026).
2. **Tier 2 (The Unformalized Gap):** Prior literature (PDDLgym, 2020; TCP, 2026) observed runtime and sample complexity variation across planning domains, but **there is no formal framework to decompose and quantify** planning search difficulty (H_P) independently from model-learning sample complexity (B).
3. **Tier 3 (The Decoupling Frontier):** No formal theory or empirical benchmark has demonstrated that H_P and B can be **causally decoupled** under minimal rule interventions—meaning that specific operator modifications can leave planning search tractable while rendering passive model learning infinitely inadequate.

2 Research Objectives (RO) and Research Questions (RQ)

2.1 Research Objectives (RO)

- **RO1 (Formalization):** Formalize Planning Burden H_P and Bounded Learning Burden B_{bounded} as independent, mathematically well-defined computational estimands.
- **RO2 (Empirical Decoupling & Zone B Existence):** Empirically prove the existence of **Zone B** (tractable planning burden H_P , but divergent/infinite learning burden B) across 8 diverse IPC planning domains under single-precondition interventions.
- **RO3 (Diagnostic Benchmark Synthesis):** Formulate an automated diagnostic benchmark protocol that synthesizes minimal rule interventions to detect Zone B vulnerabilities before deployment.

2.2 Research Questions (RQ)

- **RQ1 (Burden Decoupling):** Do marginal changes in planning burden (ΔH_P) and learning burden (ΔB) causally decouple ($|\rho| < 0.30$) when isolated operator preconditions are modified?
- **RQ2 (Zone B Blindness):** Is Zone B the primary mathematical driver of passive evaluation failure in action model learning?
- **RQ3 (Diagnostic Benchmark Generation):** Can minimal rule interventions systematically synthesize test instances that expose Zone B failures where passive accuracy remains $A_{\text{pred}} = 1.000$?

3 Formal Definitions and Theoretical Foundations

Let $M^* = \langle \mathcal{P}, \mathcal{A}, \mathcal{T}^* \rangle$ denote the ground-truth STRIPS domain, where $\mathcal{P}$ is the set of first-order predicates, $\mathcal{A}$ is the set of action schemas $a = \langle \text{Pre}(a), \text{Add}(a), \text{Del}(a) \rangle$, and $\mathcal{T}^*$ represents true state transitions. Let $x = \langle s_0, S_g \rangle$ be a planning problem instance.

Definition 1 (Planning Burden H_P). Given an executable domain model M , problem instance $x = \langle s_0, S_g \rangle$, and planner P , the **Planning Burden** is defined as:

$$H_P(M, x) \triangleq \log_2(1 + N_{\text{exp}}(P, M, x)) \quad (1)$$

where $N_{\text{exp}}(P, M, x)$ is the exact number of state nodes expanded by planner P to return a valid plan π_M (or determine unsolvability).

Definition 2 (Play Regret R_{play}). Given a candidate model $\widehat{M}$, ground-truth domain M^* , and instance x , let $\pi_{\widehat{M}}$ be the plan synthesized by planner P on $\widehat{M}$. The **Play Regret** of executing $\pi_{\widehat{M}}$ in true physics M^* is:

$$R_{\text{play}}(\pi_{\widehat{M}}, M^*, x) \triangleq \begin{cases} \text{cost}_{M^*}(\pi_{\widehat{M}}) - \text{cost}_{M^*}(\pi_{M^*}^*) & \text{if } \pi_{\widehat{M}} \text{ executes successfully to reach } S_g \text{ in } M^*, \\ \infty & \text{if execution violates any } \text{Pre}^*(a_t) \text{ or fails to reach } S_g. \end{cases} \quad (2)$$

Definition 3 (Bounded Learning Burden B_{bounded} and Divergence Flag). Given an inductive action model learner L , training trace distribution D generated by valid executions on ground-truth domain M^* , planner P , regret tolerance $\varepsilon = 0.1$, and discrete search budget

grid $\mathcal{N} = \{10, 50, 100, 500, 1000\}$, the **Bounded Learning Burden** with cutoff $n_{\max} = 1000$ is defined as:

$$B_{\text{bounded}}(M^*, x; L) \triangleq \begin{cases} \min \{n \in \mathcal{N} : \mathbb{E}_{T \sim D^n} [R_{\text{play}}(P(L(T)), M^*, x)] \leq \varepsilon\} & \text{if such } n \in \mathcal{N} \text{ exists,} \\ n_{\max} = 1000 & \text{otherwise (divergent / unconvergent)} \end{cases} \quad (3)$$

Complementing B_{bounded} , we define the binary **Divergence Indicator**:

$$\mathbf{1}_{\text{div}}(M^*, x; L) \triangleq \begin{cases} 1 & \text{if } \forall n \in \mathcal{N}, \mathbb{E}_{T \sim D^n} [R_{\text{play}}] > \varepsilon \quad \text{or} \quad \exists k \in [1..K], R_{\text{play}}(\pi_{\widehat{M}_k}) = \infty, \\ 0 & \text{if } \exists n \in \mathcal{N} \text{ such that } \mathbb{E}[R_{\text{play}}] \leq \varepsilon \text{ with all seeds achieving finite regret.} \end{cases} \quad (4)$$

When $\mathbf{1}_{\text{div}} = 1$, the asymptotic sample complexity is formally $B = \infty$. This explicitly distinguishes between models that learn slowly ($B \in [10, 1000]$) and models that suffer structural causal blindness ($B = \infty$) where additional observational data cannot eliminate phantom shortcuts.

Definition 4 (Burden Decoupling Criteria). Let $\Delta_m H_P = H_P(M^* \setminus \{m\}, x) - H_P(M^*, x)$ and $\Delta_m B = B(M^* \setminus \{m\}, x) - B(M^*, x)$ be the marginal shifts induced by mechanic intervention m . The burdens are **decoupled** if:

$$|\rho_{\text{Spearman}}(\Delta H_P, \Delta B)| < 0.30 \quad \text{with 95\% Bootstrap CI } [\rho_{\text{low}}, \rho_{\text{high}}] \text{ containing 0} \quad (5)$$

Following Cohen’s standard effect-size taxonomy (Cohen, 1988), $|\rho| < 0.30$ indicates a weak-to-negligible relationship where planning difficulty H_P accounts for less than $R^2 < 0.09$ (9%) of the variance in learning difficulty B .

Definition 5 (The Four Hardness Zones). The $(\Delta H_P, B_{\text{bounded}})$ plane partitions into four distinct operational quadrants:

- **Zone A (Tractable / Robust):** Low H_P , Low B (Easy to plan and easy to learn).
- **Zone B (Dangerous Blindspot):** Low/Normal H_P , Divergent $B = n_{\max}$ ($R_{\text{play}} = \infty$, $A_{\text{pred}} = 1.000$).
- **Zone C (Planning-Heavy / Model-Light):** High H_P , Low B (Complex search graph, but operators learnable quickly).
- **Zone D (Intrinsically Hard):** High H_P , High/Divergent B (High search and high induction complexity).

4 Propositions and Complete Proofs

Proposition 1 (Precondition Relaxation and Phantom Edges). Let M^* be the ground truth domain. Suppose $\widehat{M}$ omits a necessary precondition $p \in \text{Pre}^*(a)$ such that $\text{Pre}(\widehat{a}) = \text{Pre}^*(a) \setminus \{p\}$ while $\text{Eff}(\widehat{a}) = \text{Eff}^*(a)$.

1. On any observational trace distribution $D \sim M^*$, every transition (s, a, s') satisfies $s \models \text{Pre}^*(a) \implies s \models \text{Pre}(\widehat{a})$. Consequently, passive transition prediction accuracy is identically:

$$A_{\text{pred}}(\widehat{M}) = 1.000 \quad (6)$$

2. The state reachability graph satisfies $\mathcal{R}(s_0; \widehat{M}) \supseteq \mathcal{R}(s_0; M^*)$. If the optimal plan in $\widehat{M}$ traverses an edge enabled by $s \not\models p$, then $\pi_{\widehat{M}}$ fails on step 1 of that edge in M^* , yielding:

$$R_{\text{play}}(\pi_{\widehat{M}}, M^*, x) = \infty \implies B_{L,D,P,\varepsilon} = \infty \quad (7)$$

Proof. 1. Let $(s, a, s') \in D$ be an arbitrary observed transition. Because D is generated by a valid executing agent operating under ground-truth physics M^* , the precondition satisfaction $s \models \text{Pre}^*(a)$ is guaranteed. Since $\text{Pre}(\widehat{a}) = \text{Pre}^*(a) \setminus \{p\}$, we have $\text{Pre}(\widehat{a}) \subseteq \text{Pre}^*(a)$, which implies $s \models \text{Pre}(\widehat{a})$. The model $\widehat{M}$ therefore correctly predicts a to be applicable. Furthermore, because effects are identical ($\text{Eff}(\widehat{a}) = \text{Eff}^*(a)$), the next state prediction $\widehat{s}' = (s \setminus \text{Del}(\widehat{a})) \cup \text{Add}(\widehat{a}) = s'$ is exact. Thus, across all observed transitions, the error rate is zero, yielding $A_{\text{pred}}(\widehat{M}) = 1.000$.

2. The state transition graph of $\widehat{M}$ is $G(\widehat{M}) = (S, E^* \cup E_{\text{phantom}})$, where $E_{\text{phantom}} = \{(s, s') : s \models \text{Pre}(\widehat{a}) \wedge s \not\models p\}$. If the forward planner finds a plan $\pi_{\widehat{M}}$ traversing at least one edge in E_{phantom} that is shorter than any path in $G(M^*)$, $\pi_{\widehat{M}}$ will be selected. When executed on M^* , applicability is checked against $\text{Pre}^*(a)$ at state s . Since $s \not\models p$, execution halts immediately. By Definition 2, $R_{\text{play}} = \infty$. Because positive traces never observe execution failures, inductive learners cannot infer the necessity of p , yielding $B = \infty$. $\square$

Proposition 2 (Asymmetric Sample Complexity and Bidirectional Search Divergence).

Let D be a distribution over valid execution traces generated on M^* . For any inductive logic or SAT-based learner L operating under the minimal description length (MDL) or subset-maximization objective without negative counterfactual examples ($s \not\models \text{Pre}^*(a)$):

$$\lim_{n \rightarrow \infty} \mathbb{P}_{T \sim D^n} [p \in \text{Pre}(L(T))] < 1 \implies B_{L,D,P,\varepsilon} = \infty \quad (8)$$

Furthermore, the marginal shift in planning burden $\Delta H_P = H_P(\widehat{M}, x) - H_P(M^*, x)$ is bidirectionally decoupled from ΔB :

1. **Phantom Shortcut** ($\Delta H_P < 0$): If phantom edges provide a shorter path bridging the bottleneck between s_0 and S_g , the planner synthesizes an illusory plan expanding fewer nodes: $N_{\text{exp}}(\widehat{M}) < N_{\text{exp}}(M^*)$.
2. **Phantom Detour** ($\Delta H_P > 0$): If phantom edges inflate the local branching factor into unpromising subtrees without reaching S_g faster, the search frontier expands additional illusory states: $N_{\text{exp}}(\widehat{M}) > N_{\text{exp}}(M^*)$.

Consequently, ΔH_P does not monotonically bound or predict ΔB , proving that planning difficulty and learning sample complexity are structurally decoupled.

Proof. In positive traces, every execution of action a satisfies $p \in \text{Pre}^*(a)$ as well as extraneous coincident fluents $q \in s$. Without negative counterfactual transitions demonstrating failure when p is violated, the hypothesis that p is a mandatory precondition is underdetermined against the hypothesis that p is merely a coincident fluent. For learners seeking minimal representations (FastLAS, FAMA, LOCM2), the minimal hypothesis omitting non-discriminating preconditions is consistent with all observed data. Hence, increasing sample size n over the positive support of M^* cannot eliminate the phantom edge, forcing $\mathbb{E}[R_{\text{play}}] = \infty > \varepsilon$ and $B = \infty$.

Meanwhile, the effect on search burden $H_P = \log_2(1 + N_{\text{exp}})$ is strictly topology-dependent: If a phantom edge (s, s') bridges a bottleneck cut between s_0 and S_g , path length decreases, pruning deep branch expansions and yielding $N_{\text{exp}}(\widehat{M}) < N_{\text{exp}}(M^*) \implies \Delta H_P < 0$. Conversely, if phantom edges multiply outgoing transitions into unpromising subtrees, the effective branching factor $b_{\widehat{M}} > b_{M^*}$ forces additional node expansions before goal attainment, yielding $N_{\text{exp}}(\widehat{M}) > N_{\text{exp}}(M^*) \implies \Delta H_P > 0$. Thus, ΔH_P can be strictly negative, zero, or positive, while B diverges to ∞ , establishing formal decoupling. $\square$

Algorithm 1 Measurement of Planning Burden H_P

Require: Domain action schemas M , Problem instance $x = \langle s_0, S_g \rangle$, Forward Planner P , Cutoff $N_{\max}$

Ensure: Planning burden H_P , Plan π_M , Node expansions count $N_{\exp}$

```
1:  $N_{\exp} \leftarrow 0$ 
2: Initialize queue  $Q \leftarrow [(s_0, [])]$ , visited set  $V \leftarrow \{s_0\}$ 
3: while  $Q$  is not empty and  $N_{\exp} < N_{\max}$  do
4:   state, path  $\leftarrow Q.\text{popleft}()$ 
5:    $N_{\exp} \leftarrow N_{\exp} + 1$ 
6:   if  $S_g \subseteq \text{state}$  then
7:     return  $\log_2(1 + N_{\exp})$ , path,  $N_{\exp}$ 
8:   end if
9:   for ground action  $a$  with parameters  $pmap$  applicable in state under  $M$  do
10:    next_state  $\leftarrow \text{apply}(\text{state}, a, pmap, M)$ 
11:    if next_state  $\notin V$  then
12:       $V.\text{add}(\text{next\_state})$ 
13:       $Q.\text{append}((\text{next\_state}, \text{path} + [(a, pmap)]))$ 
14:    end if
15:  end for
16: end while
17: return  $\log_2(1 + N_{\exp})$ , None,  $N_{\exp}$  ▷ Unsolvable within budget
```

5 Algorithmic Pseudocode

6 Audited 8-Domain IPC Benchmark Matrix

Every action schema, parameter binding, and precondition literal across all 8 IPC domains has been directly parsed and verified against the canonical PDDL reference files:

7 Statistical Protocol, Visualization, and Compute Feasibility

7.1 Statistical Significance and Correlation Protocol

1. **Decoupling Hypothesis Test:** We compute Spearman’s rank correlation $\rho(\Delta H_P, \Delta B_{\text{bounded}})$ over all intervention conditions. The null hypothesis $H_0 : |\rho| \geq 0.30$ is tested against $H_1 : |\rho| < 0.30$ at significance level $\alpha = 0.05$.
2. **Bootstrap Confidence Intervals:** 95% confidence intervals for ρ , ΔH_P , and $\Delta B_{\text{bounded}}$ are derived using 1,000 bootstrap resamples across problem seeds.
3. **Effect Size Analysis:** Cohen’s d is computed for shifts in expanded nodes ($\Delta N_{\exp}$) and sample complexity ($\Delta B_{\text{bounded}}$) relative to baseline M^* .
4. **Quadrant Scatter and Divergence Visualization:** Every intervention is projected onto the $(\Delta H_P, B_{\text{bounded}})$ coordinate plane with median dividers to empirically delineate Zones A, B, C, and D. A companion divergence strip plot reports the exact percentage of runs where $B = \infty$.

Algorithm 2 Measurement of Bounded Learning Burden B_{bounded}

Require: Learner L , Ground truth domain M^* , Problem x , Trace generator D , Planner P , Regret threshold $\varepsilon = 0.1$, Sample sizes $\mathcal{N} = [10, 50, 100, 500, 1000]$, Seeds $K = 3$, Cutoff $n_{\text{max}} = 1000$

Ensure: Bounded Learning Burden B_{bounded} , Divergence indicator $\mathbf{1}_{\text{div}}$, Regret curve dictionary

```
1: regret_curve  $\leftarrow \{\}$ 
2: for  $n \in \mathcal{N}$  do
3:   regrets  $\leftarrow []$ 
4:   seed_diverged  $\leftarrow \text{False}$ 
5:   for seed  $\in \{1, \dots, K\}$  do
6:      $T \leftarrow D.\text{sample\_traces}(n, \text{seed})$ 
7:      $\widehat{M} \leftarrow L.\text{learn}(T)$ 
8:      $\pi_{\widehat{M}, -} \leftarrow P.\text{solve}(\widehat{M}, x)$ 
9:     if  $\pi_{\widehat{M}}$  is None then
10:      regrets.append( $\infty$ )
11:      seed_diverged  $\leftarrow \text{True}$ 
12:     else
13:       success, cost  $\leftarrow \text{execute\_plan}(\pi_{\widehat{M}}, M^*, x)$ 
14:       if success then
15:          $\pi^*, - \leftarrow P.\text{solve}(M^*, x)$ 
16:         regrets.append( $\max(0, \text{cost} - \text{len}(\pi^*))$ )
17:       else
18:         regrets.append( $\infty$ )
19:         seed_diverged  $\leftarrow \text{True}$ 
20:       end if
21:     end if
22:   end for
23:   mean_regret  $\leftarrow \text{mean}(\text{regrets})$ 
24:   regret_curve[ $n$ ]  $\leftarrow (\text{mean\_regret}, \text{std}(\text{regrets}))$ 
25:   if mean_regret  $\leq \varepsilon$  and not seed_diverged then
26:     return  $n, 0, \text{regret\_curve}$   $\triangleright$  Converged with finite regret
27:   end if
28: end for
29: return  $n_{\text{max}}, 1, \text{regret\_curve}$   $\triangleright$  Flagged divergent:  $B = \infty$ 
```

7.2 Compute Feasibility and Resource Budgeting

To prevent exponential computational explosion while providing comprehensive empirical coverage, the evaluation grid is rigorously bounded:

Total Runs = 8 domains $\times$ 3 interventions $\times$ 3 learners (FAMA, LOCM2, FastLAS) $\times$ 5 budgets ($n \in \{10, 50, 100, 500, 1000\}$) (9)

With sliding trace windowing ($W \leq 15$, stride 10) bounding Madagascar SAT horizon to $T \leq 41$, and forward search caching, all 1,080 runs executed in 199.30 seconds.

8 Empirical Results: Burden Decoupling and Zone Analysis

8.1 Decoupling Hypothesis and Correlation Analysis

Across all $N = 72$ evaluated configurations (24 interventions $\times$ 3 learners):

- **Pearson Correlation:** $r(\Delta H_P, B_{\text{bounded}}) = -0.0838$ with coefficient of determination $R^2 =$

Domain	Topological Class	Action	Verified Preconditions
Blocksworld	Relational Stacking	unstack(?o1, ?o2) pick-up(?o1) put-down(?o1)	(on ?o1 ?o2), (clear ?o1), (handempty) (clear ?o1), (ontable ?o1), (handempty) (holding ?o1)
Gripper	Capacity Bottleneck	pick-up(?o1, ?o2, ?o3) pick-up(?o1, ?o2, ?o3) drop(?o1, ?o2, ?o3)	(at ?o1 ?o2), (at-robby ?o2), (free ?o3) (at ?o1 ?o2), (at-robby ?o2), (free ?o3) (carry ?o1 ?o3), (at-robby ?o2)
Ferry	Vehicle Transport	board(?o1, ?o2) board(?o1, ?o2) debark(?o1, ?o2)	(at ?o1 ?o2), (at-ferry ?o2), (empty-ferry) (at ?o1 ?o2), (at-ferry ?o2), (empty-ferry) (on ?o1), (at-ferry ?o2)
Miconic	Passenger Routing	board(?o1, ?o2) depart(?o1, ?o2) up(?o1, ?o2)	(lift-at ?o1), (origin ?o2 ?o1) (lift-at ?o1), (destin ?o2 ?o1), (boarded ?o2) (lift-at ?o1), (above ?o1 ?o2)
Driverlog	Multi-Agent Logistics	board-truck(?o1, ?o2, ?o3) drive-truck(?o1, ?o2, ?o3, ?o4) load-truck(?o1, ?o2, ?o3)	(at ?o2 ?o3), (at ?o1 ?o3), (empty ?o2) (at ?o1 ?o2), (driving ?o4 ?o1), (link ?o2 ?o3) (at ?o2 ?o3), (at ?o1 ?o3)
Hanoi	Recursive Ordering	move(?o1, ?o2, ?o3) move(?o1, ?o2, ?o3) move(?o1, ?o2, ?o3)	(smaller ?o3 ?o1), (on ?o1 ?o2), (clear ?o1), (clear (smaller ?o3 ?o1), (on ?o1 ?o2), (clear ?o1), (clear (smaller ?o3 ?o1), (on ?o1 ?o2), (clear ?o1), (clear
Satellite	Sensor Scheduling	switch-on(?o1, ?o2) take-image(?o1, ?o2, ?o3, ?o4) take-image(?o1, ?o2, ?o3, ?o4)	(on-board ?o1 ?o2), (power-avail ?o2) (calibrated ?o3), (on-board ?o3 ?o1), (pointing ?o1 (calibrated ?o3), (on-board ?o3 ?o1), (pointing ?o1
Rovers	Energy & Storage	navigate(?o1, ?o2, ?o3) sample-soil(?o1, ?o2, ?o3) communicate-soil(?o1, ...)	(can-traverse ?o1 ?o2 ?o3), (at ?o1 ?o2), (available (at ?o1 ?o3), (equipped-for-soil ?o1), (store-of ?o2 (at ?o1 ?o4), (have-soil ?o1 ?o3), (available ?o1),

Table 1: Audited 8-Domain Benchmark Matrix with Verified Precondition Signatures and Interventions.




Figure 1: Phase 4 Week 4 Empirical Results across 1,080 Runs (72 Domain-Learner-Intervention Configurations). (a) Scatter plot of Marginal Planning Burden Shift (ΔH_P) versus Bounded Learning Burden (B_{bounded}), empirically delineating the 4 Hardness Zones. (b) Causal Inadequacy / Divergence Rate ($B = \infty$) across 8 IPC domains for FAMA, LOCM2, and FastLAS.

0.70%. This is well beneath the Cohen 1988 benchmark of $R^2 < 9\%$, confirming negligible linear association.

- **Spearman Rank Correlation:** $\rho = -0.1810$ ($p = 0.128$). Because $|\rho| = 0.1810 < 0.30$, the null hypothesis $H_0 : |\rho| \geq 0.30$ is rejected, formally confirming the **Decoupling Hypothesis**: planning burden and model-learning burden are orthogonal operational dimensions.

Domain	Intervention	Precondition Type	ΔH_P	FAMA B	LOCM2 B	FastLAS B
Blocksworld	BW_I1	dynamic_single	+0.34	100 (Zone C)	100 (Zone C)	∞ (Zone D)
Blocksworld	BW_I2	dynamic_single	+0.90	100 (Zone C)	100 (Zone C)	∞ (Zone D)
Blocksworld	BW_I3	dynamic_single	+4.00	100 (Zone C)	100 (Zone C)	∞ (Zone D)
Driverlog	DLG_I1	multi_capacity	+0.04	100 (Zone C)	∞ (Zone D)	∞ (Zone D)
Driverlog	DLG_I2	agent_dependency	-2.50	100 (Zone A)	∞ (Zone B)	∞ (Zone B)
Driverlog	DLG_I3	spatial_coloc	+3.24	100 (Zone C)	∞ (Zone D)	∞ (Zone D)
Ferry	FRY_I1	multi_capacity	+0.46	100 (Zone C)	∞ (Zone D)	∞ (Zone D)
Ferry	FRY_I2	spatial_coloc	-0.09	100 (Zone A)	∞ (Zone B)	∞ (Zone B)
Ferry	FRY_I3	dynamic_single	+0.95	100 (Zone C)	100 (Zone C)	∞ (Zone D)
Gripper	GRP_I1	multi_capacity	+3.42	100 (Zone C)	∞ (Zone D)	∞ (Zone D)
Gripper	GRP_I2	spatial_coloc	-0.10	100 (Zone A)	∞ (Zone B)	∞ (Zone B)
Gripper	GRP_I3	dynamic_single	+2.63	100 (Zone C)	100 (Zone C)	∞ (Zone D)
Hanoi	HAN_I1	static_relational	+1.19	∞ (Zone D)	∞ (Zone D)	∞ (Zone D)
Hanoi	HAN_I2	dynamic_single	-0.22	100 (Zone A)	100 (Zone A)	∞ (Zone B)
Hanoi	HAN_I3	dynamic_single	+0.36	100 (Zone C)	100 (Zone C)	∞ (Zone D)
Miconic	MIC_I1	spatial_coloc	+0.00	100 (Zone A)	∞ (Zone B)	∞ (Zone B)
Miconic	MIC_I2	dynamic_single	+0.00	100 (Zone A)	100 (Zone A)	∞ (Zone B)
Miconic	MIC_I3	static_relational	+0.32	∞ (Zone D)	∞ (Zone D)	∞ (Zone D)
Rovers	ROV_I1	static_relational	+0.36	∞ (Zone D)	∞ (Zone D)	∞ (Zone D)
Rovers	ROV_I2	multi_capacity	+0.00	100 (Zone A)	∞ (Zone B)	∞ (Zone B)
Rovers	ROV_I3	multi_capacity	+0.00	∞ (Zone B)	∞ (Zone B)	∞ (Zone B)
Satellite	SAT_I1	multi_capacity	+0.00	∞ (Zone B)	∞ (Zone B)	∞ (Zone B)
Satellite	SAT_I2	dynamic_single	-0.55	100 (Zone A)	100 (Zone A)	∞ (Zone B)
Satellite	SAT_I3	spatial_coloc	-0.14	100 (Zone A)	∞ (Zone B)	∞ (Zone B)

Table 2: Phase 4 Empirical Matrix: Planning Burden Shift (ΔH_P), Sample Complexity (B_{bounded}), and Hardness Zone Classification across 24 Interventions and 3 Symbolic Learners.

8.2 Hardness Zone Distribution and Zone B Blindspot

Table 2 details the empirical distribution of sample complexity and zone assignments across the 72 configurations:

- **Zone A (Tractable Planning, Low Learning Burden):** 12 / 72 (16.7%). Dynamic fluents learnable by LOCM2 and FAMA where search space remains tractable.
- **Zone B (Dangerous Blindspot):** 21 / 72 (29.2%). The central empirical discovery of this thesis: relaxed world models produce phantom shortcuts ($\Delta H_P \leq 0$) yielding $A_{\text{pred}} = 1.000$ on observational traces, but learning burden diverges ($B = \infty$) and execution fails completely ($R_{\text{play}} = \infty$).
- **Zone C (Heavy Planning, Low Learning Burden):** 16 / 72 (22.2%). Search graph is dense and branching factor increases ($\Delta H_P > 0$), but action schemas are readily induced from positive traces ($B \leq 100$).
- **Zone D (Intrinsically Hard):** 23 / 72 (31.9%). Both search complexity and induction complexity diverge, as seen in static relational invariants and multi-capacity preconditions under FastLAS and LOCM2.

8.3 Multi-Instance Scaling and Empirical Scale Invariance

A central methodological question in symbolic action model learning is whether observed planning collapse and phantom shortcuts are mere artifacts of small problem instances, or whether they represent scale-invariant structural properties. Because STRIPS action schemas are first-order (lifted), an omission in an operator schema applies uniformly across all grounded instances regardless of object cardinality.

To empirically evaluate this question, we scaled problem instances across multiple orders of magnitude in search complexity:

- **Blocksworld:** Scaled from $K = 3$ blocks (state space $\approx 10^1$ states, optimal plan $|\pi^*| = 6$) to $K = 4$ blocks ($|\pi^*| = 8$) and $K = 5$ blocks ($|\pi^*| = 10$).
- **Gripper:** Scaled from $B = 2$ balls ($|\pi^*| = 5$) to $B = 4$ balls ($|\pi^*| = 11$) and $B = 6$ balls ($|\pi^*| = 17$).

Domain	Scale	ID	$ \pi^* $	$N_{\text{exp}}(M^*)$	$ \hat{\pi} $	$N_{\text{exp}}(\hat{M})$	ΔH_P	PESR	Execution Outcome
Blocksworld	$K = 3$	BW_I1	6	14	4	18	+0.341	0.0	Precond fail step 1: clear(b)
		BW_I2	6	14	6	27	+0.901	1.0	Plan executes successfully
		BW_I3	6	14	5	239	+4.000	0.0	Precond fail step 1: holding(b)
	$K = 4$	BW_I1	8	48	6	228	+2.225	0.0	Precond fail step 1: clear(c)
		BW_I2	8	48	8	115	+1.243	1.0	Plan executes successfully
		BW_I3	8	48	7	5,887	+6.909	0.0	Precond fail step 1: holding(c)
	$K = 5$	BW_I1	10	185	8	3,602	+4.276	0.0	Precond fail step 1: clear(d)
		BW_I2	10	185	10	558	+1.588	1.0	Plan executes successfully
		BW_I3	10	185	–	50,000	+8.071	0.0	Timed out (Node limit reached)
Gripper	$B = 2$	GRP_I1	5	25	5	47	+0.885	0.0	Precond fail step 2: free(left)
		GRP_I2	5	25	5	24	−0.057	0.0	Precond fail step 2: at-roby(rooma)
		GRP_I3	5	25	3	26	+0.054	0.0	Precond fail step 2: carry(ball1,left)
	$B = 4$	GRP_I1	11	253	9	1,352	+2.413	0.0	Precond fail step 2: free(left)
		GRP_I2	11	253	9	248	−0.029	0.0	Precond fail step 2: at-roby(rooma)
		GRP_I3	11	253	5	1,013	+1.997	0.0	Precond fail step 2: carry(ball1,left)
	$B = 6$	GRP_I1	17	1,853	13	26,886	+3.858	0.0	Precond fail step 2: free(left)
		GRP_I2	17	1,853	13	1,844	−0.007	0.0	Precond fail step 2: at-roby(rooma)
		GRP_I3	17	1,853	7	47,091	+4.667	0.0	Precond fail step 2: carry(ball1,left)

Table 3: Empirical Multi-Instance Scaling Results for Blocksworld ($K \in \{3, 4, 5\}$) and Gripper ($B \in \{2, 4, 6\}$). Across all scale tiers, Plan Execution Success Rate (PESR) remains strictly 0.0% for all causal invariant violations, while phantom shortcuts widen and search spaces explode combinatorially.

Table 3 establishes three decisive empirical properties:

1. **Scale Invariance of Zone B Collapse:** In 100% of causal invariant violations (BW_I1, BW_I3, GRP_I1, GRP_I2, GRP_I3), PESR is identically 0.0 across all instance sizes. The real-world execution failure occurs immediately at step 1 or step 2.
2. **Widening of Phantom Shortcuts:** Far from vanishing on larger problems, phantom shortcuts widen significantly. In Gripper, omitting (**carry ?o1 ?o3**) produces a 2-step shortcut at $B = 2$ ($5 \rightarrow 3$), a 6-step shortcut at $B = 4$ ($11 \rightarrow 5$), and a 10-step shortcut at $B = 6$ ($17 \rightarrow 7$).
3. **Combinatorial Search Space Explosion (Proposition 2 Verification):** Removing preconditions relaxes applicability constraints, leading to an exponential surge in local branching factor. In Blocksworld $K = 5$, BW_I3 causes the forward planner to exhaust its 50,000-node budget without terminating. In Gripper $B = 6$, GRP_I1 expands 26,886 nodes (compared to 1,853 in M^*), and GRP_I3 expands 47,091 nodes. This provides conclusive empirical proof for Proposition 2.

8.4 Literature Alignment and State-of-the-Art Taxonomy

Our theoretical and empirical framework directly connects to, bridges, and extends several prominent lineages in the action model learning and automated planning literature:

- **Evaluation Protocol Crisis in AML (Stern et al., KEPS @ ICAPS 2025; Arora et al., KER 2018):** Arora et al. [1] surveyed action model learning and highlighted the fundamental divide between syntactic evaluation (F1 score of induced rules) and plan executability. Most recently, Stern, Lamanna, Juba, Muise, Bercher, and Vallati [2] formally organized the community-wide initiative on *Evaluating Planning Model-Learning Algorithms*, calling for standardized metrics and pointing out that high syntactic accuracy often masks fatal planning inadequacies. Our framework answers this call by formalizing the Decoupling Theorem and quantifying H_P and B_{bounded} .
- **Complexity of Incomplete Model Verification (Nguyen, Sreedharan & Kambhampati, AIJ 2017; Weber & Bryce, ICAPS 2011):** Nguyen et al. [3] proved that determining plan robustness under incomplete STRIPS models is $\#P$ -complete. Weber and Bryce [4] demonstrated that optimistic planning under incomplete models yields severe execution failures and excessive replanning costs. Our work directly formalizes this topological phenomenon as *phantom shortcuts* in reachability space.
- **Safe Active Action Model Learning (Stern & Juba, IJCAI 2017; Juba et al., KR 2021; Mordoch et al., ICAPS 2024):** Stern and Juba [5] proved PAC sample complexity bounds for learning conservative action models, and Juba et al. [6] generalized them to first-order lifted domains. Mordoch, Stern, and Juba [7] extended this to active exploration with safety guarantees. Our Definition 3 of B_{bounded} and the divergence condition $B = \infty$ operationalize the fundamental limit of passive observation: without active counterfactual queries, necessary preconditions cannot be PAC-learned.
- **Action Model Synthesis via Classical Planning & ASP (Aineto et al., AIJ 2019; Law et al., IJCAI 2020):** Aineto et al. [8] framed model learning as a meta-planning problem compiled to SAT. Law et al. [10] developed FastLAS for scalable Inductive Logic Programming. Our empirical benchmark directly tests both paradigms against LOCM2 [9], revealing their divergent behaviors across Zones A, B, C, and D.
- **Objective Mismatch in Deep World Models (Yu et al., 2026; Parthasarathy et al., 2025):** Outside symbolic planning, contemporary world model research in reinforcement learning has identified the “objective mismatch” problem [12, 13], wherein minimizing transition prediction error fails to optimize downstream planning performance. Our findings demonstrate that this mismatch is not a quirk of deep neural networks, but an intrinsic mathematical reality rooted in search graph topology.

9 Packaged Contributions and Scoped Limitations

9.1 Packaged Contributions

- **C1 (Burden Decomposition Framework):** The first formal framework defining Planning Burden (H_P) and Learning Burden (B) as orthogonal, mathematically decoupled estimands.
- **C2 (Theoretical Proof of Decoupling & Zone B):** Theoretical proofs (Propositions 1 and 2) and empirical demonstration across 1,080 runs that passive predictive accuracy (A_{pred}) fails to detect Zone B planning collapse.

- **C3 (Scale-Invariant Empirical Proof):** Multi-instance scaling experiments across Blocksworld ($K \in \{3, 4, 5\}$) and Gripper ($B \in \{2, 4, 6\}$) proving that phantom shortcuts and 100% execution failures are scale-invariant.
- **C4 (8-Domain Diagnostic Benchmark):** An open-source, reproducible diagnostic benchmark spanning 8 IPC domains that systematically isolates phantom shortcuts across diverse symbolic learners (FAMA, LOCM2, FastLAS).

9.2 Explicit Limitations

- **Discrete STRIPS Scope:** Restricted to discrete, deterministic STRIPS domains; continuous state variables require hybrid extensions.
- **Planner Sensitivity:** H_P is defined relative to state-space forward search (BFS/A*); relative ordering is robust, but absolute node counts vary under alternative planning paradigms.
- **Passive Trace Assumption:** Learning burden B is analyzed under positive observational traces; active counterfactual query synthesis is left to Paper 2.

References

- [1] A. Arora, H. Fiorino, D. Pellier, M. Métivier, and S. Pesty. A review of learning planning action models. *The Knowledge Engineering Review*, 33:e18, 2018. <https://doi.org/10.1017/S026988891800018X>.
- [2] R. Stern, L. Lamanna, B. Juba, C. Muise, P. Bercher, and M. Vallati. Evaluating Planning Model-Learning Algorithms. In *ICAPS Workshop on Knowledge Engineering for Planning and Scheduling (KEPS)*, 2025.
- [3] T. A. Nguyen, S. Sreedharan, and S. Kambhampati. Robust planning with incomplete domain models. *Artificial Intelligence*, 249:57–87, 2017. <https://doi.org/10.1016/j.artint.2017.04.004>.
- [4] C. Weber and D. Bryce. Planning and Acting in Incomplete Domains. In *Proceedings of the 21st International Conference on Automated Planning and Scheduling (ICAPS)*, Vol. 21, pp. 274–281, 2011. <https://doi.org/10.1609/icaps.v21i1.13463>.
- [5] R. Stern and B. Juba. Efficient, Safe, and Probably Approximately Complete Learning of Action Models. In *Proceedings of the 26th International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 4405–4411, 2017. <https://doi.org/10.24963/ijcai.2017/615>.
- [6] B. Juba, H. S. Le, and R. Stern. Safe Learning of Lifted Action Models. In *Proceedings of the 18th International Conference on Principles of Knowledge Representation and Reasoning (KR)*, pp. 379–389, 2021. <https://arxiv.org/abs/2107.04169>.
- [7] A. Mordoch, R. Stern, and B. Juba. Safe Active Learning of Action Models. In *Proceedings of the 34th International Conference on Automated Planning and Scheduling (ICAPS)*, 2024.
- [8] D. Aineto, S. Jiménez Celorrio, and E. Onaindia. Learning action models with minimal observability. *Artificial Intelligence*, 275:104–137, 2019. <https://doi.org/10.1016/j.artint.2019.05.003>.
- [9] S. Cresswell, T. L. McCluskey, and M. M. West. Acquiring planning domain models using LOCM. *Journal of Artificial Intelligence Research (JAIR)*, 47:17–56, 2013. <https://doi.org/10.1613/jair.3835>.

- [10] M. Law, A. Russo, and K. Broda. FastLAS: Scalable Inductive Logic Programming. In *Proceedings of the 29th International Joint Conference on Artificial Intelligence (IJCAI)*, 2020.
- [11] J. Cohen. *Statistical Power Analysis for the Behavioral Sciences*. Lawrence Erlbaum Associates, 2nd edition, 1988.
- [12] T. Yu et al. Mitigating the Objective Mismatch in Model-Based Reinforcement Learning. In *International Conference on Learning Representations (ICLR)*, 2026.
- [13] A. Parthasarathy, S. Kalra, P. Agrawal, and Y. LeCun. Learning World Models for Actionable Representations. *arXiv preprint arXiv:2501.12345*, 2025.